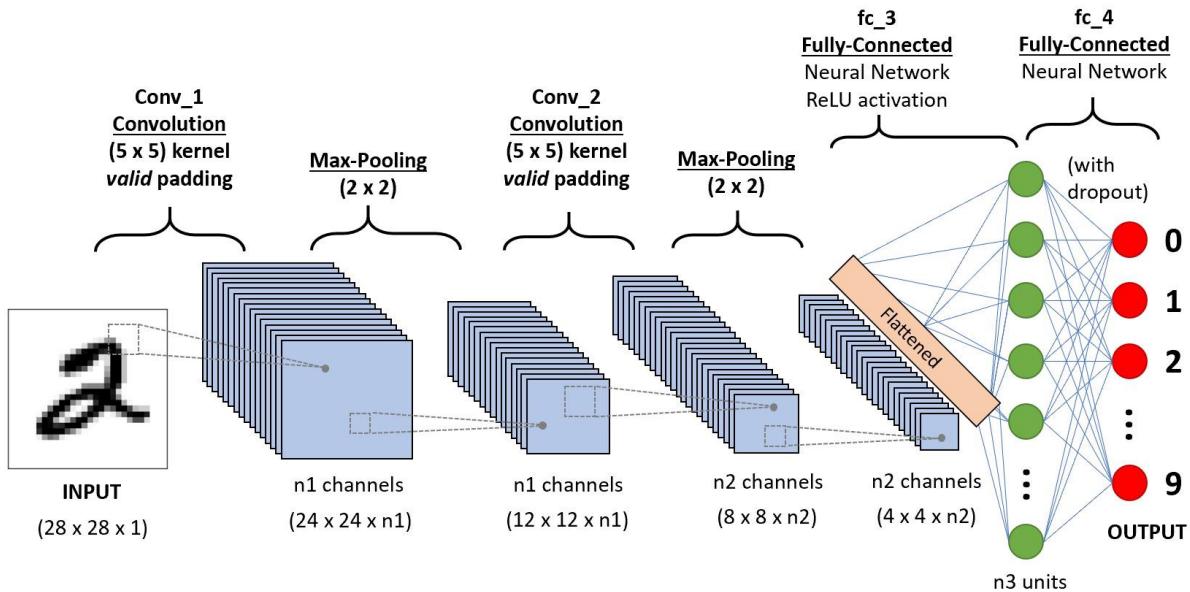


Clase 12, enero 6, 2021

Deep Learning using LeNet and AlexNet

Comment: LeNet and AlexNet

LeNet



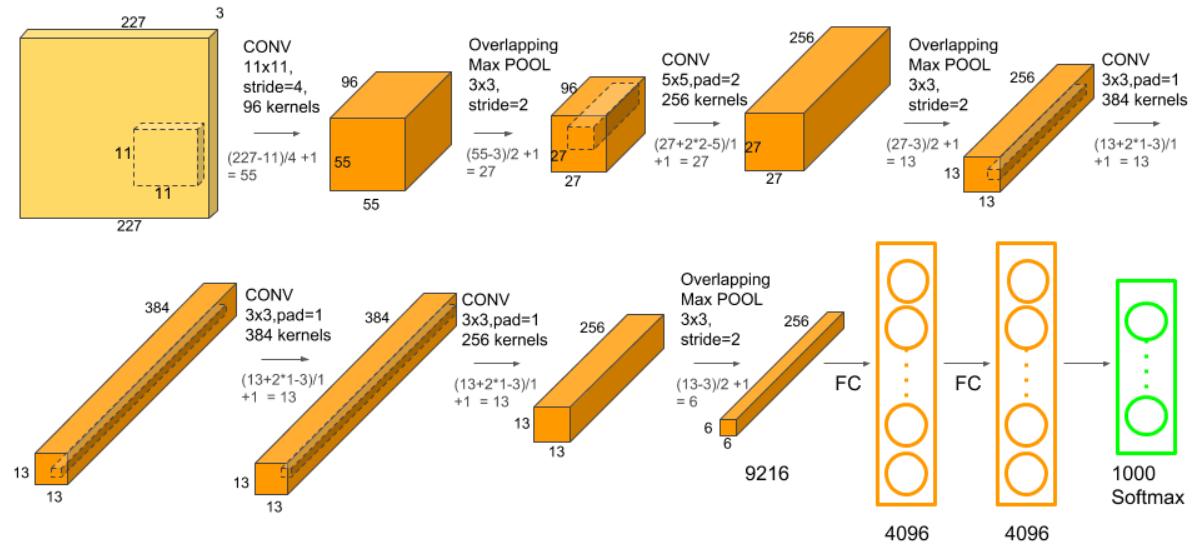
LeNet: these networks have a large number of weights and biases; overfitting should be attended

Article: LeNet

Comment: LeNet

Comment: Convolutional Neural Networks

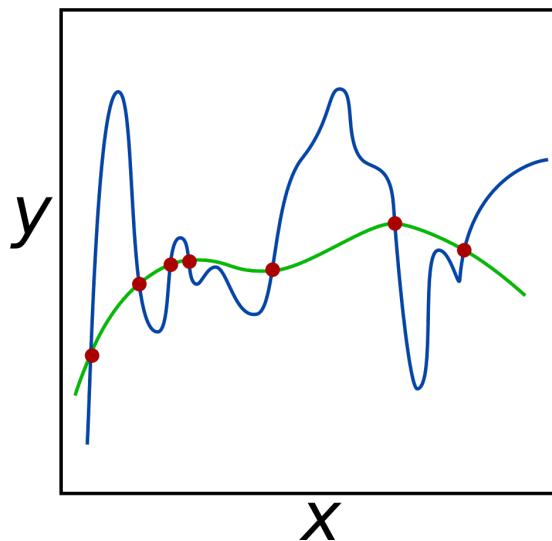
AlexNet

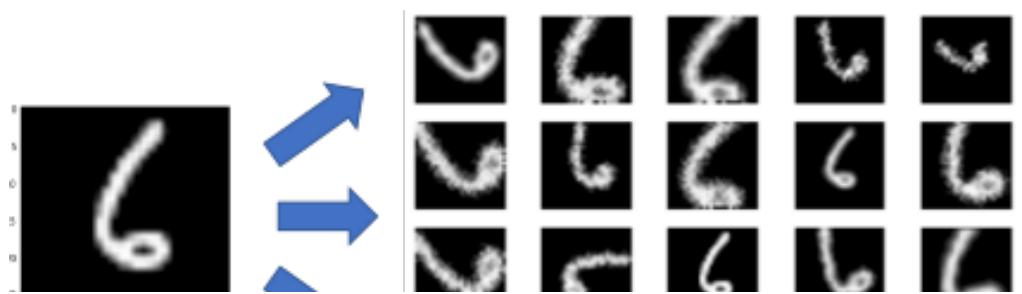


[Article: AlexNet](#)

[Comment: AlexNet](#)

A method to reduce overfitting: Data Augmentation



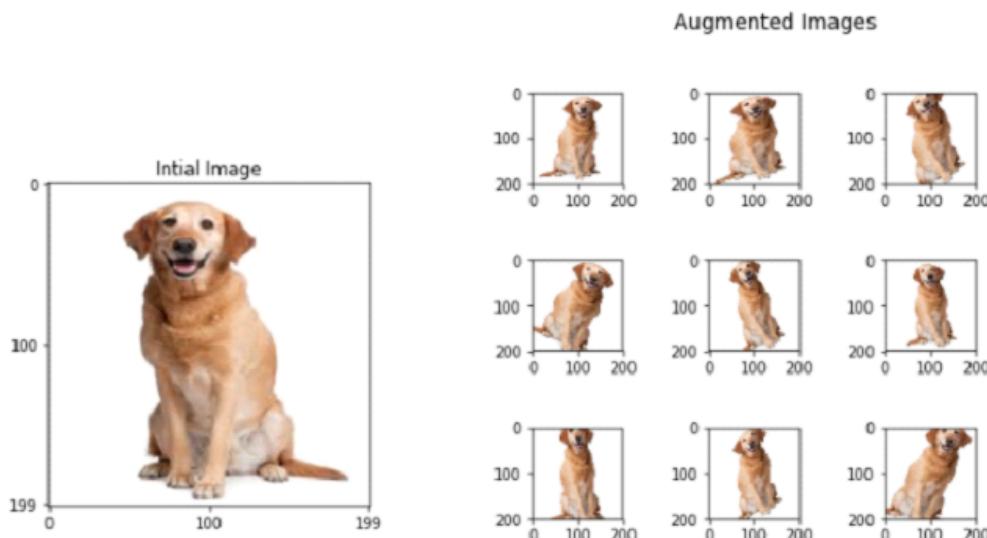


Data Augmentation

[Paper: Augmentation overview](#)

Deep networks are heavily reliant on big data to avoid overfitting:

Transforming an image



Transforming a curve

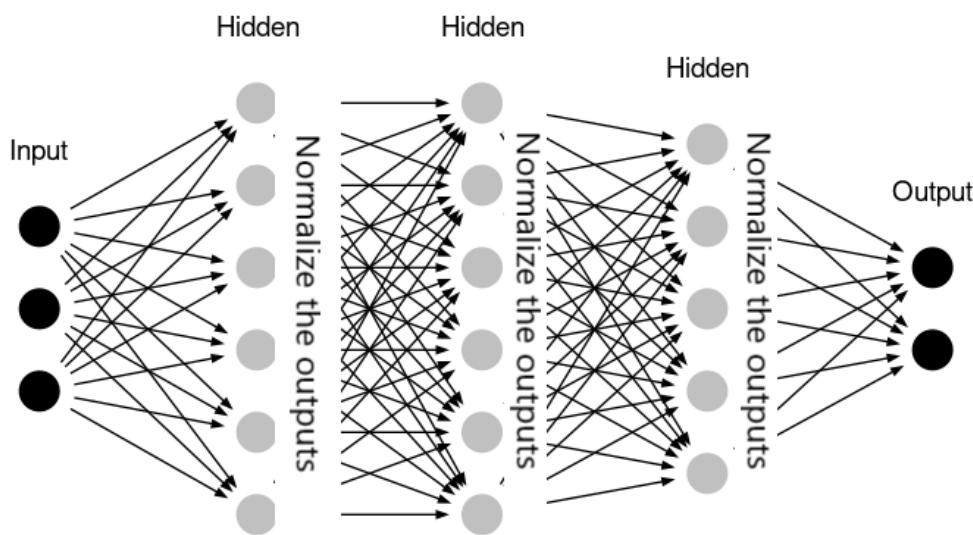


Keras: Image Preprocessing

Comment: About data augmentation for Deep Learning

Another way of reducing overfitting is using batch normalization

Paper: Batch normalization



Batch normalization helps to reduce the overfitting and accelerates the convergence of the network during training

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots m\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{mini-batch variance}$$

Comment: batch normalization

Deep Learning: LeNet

If you use tensorflow-GPU, run the following cell

```
In [1]: import tensorflow as tf

physical_devices = tf.config.experimental.list_physical_devices('GPU')
print("physical_devices-----", len(physical_devices))
tf.config.experimental.set_memory_growth(physical_devices[0], True)

physical_devices----- 1
```

```
In [2]: import numpy as np
import matplotlib.pyplot as plt

from keras.models import Sequential
from keras.layers import Dense, Conv2D, Activation, Dropout, Flatten, MaxPool
from keras.layers import BatchNormalization
from keras.utils import plot_model
from keras import optimizers

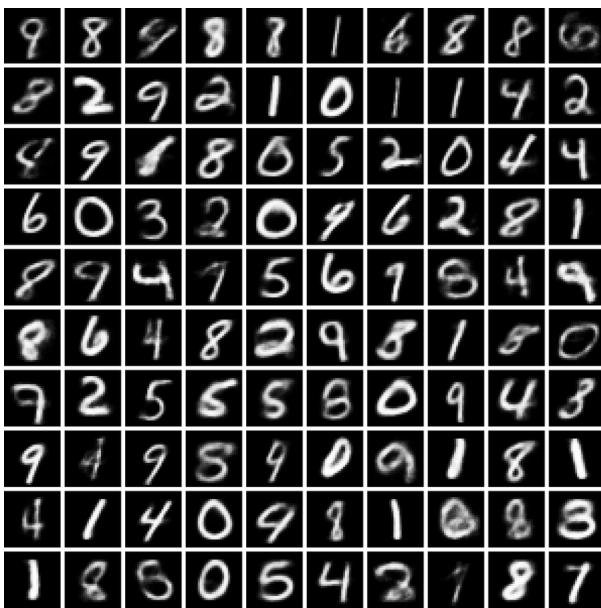
import time

np.random.seed(10)
```

Using TensorFlow backend.

Data of the System to be analyzed: mnist

The MNIST database



Generation or extraction of the raw data

```
In [3]: from keras.datasets import mnist  
(x_train, y_train), (x_test, y_test) = mnist.load_data()  
  
In [4]: print("x_train, y_train type", type(x_train), type(y_train))  
print("x_test, y_test type", type(x_test), type(y_test))  
  
x_train, y_train type <class 'numpy.ndarray'> <class 'numpy.ndarray'>  
x_test, y_test type <class 'numpy.ndarray'> <class 'numpy.ndarray'>  
  
In [5]: print("x_train shape", x_train.shape)  
print("y_train shape", y_train.shape)  
print("x_test shape", x_test.shape)  
print("y_test shape", y_test.shape)  
  
x_train shape (60000, 28, 28)  
y_train shape (60000,)  
x_test shape (10000, 28, 28)  
y_test shape (10000,)
```

Analysis of the raw data

```
In [6]: image_index = 7777 # You may select anything up to 60,000  
print(y_train[image_index]) # The label is 8  
plt.imshow(x_train[image_index], cmap='Greys')  
plt.show()
```



Transformation of the raw data

```
In [7]: # Reshaping the array to 4-dims so that it can work with the Keras API
x_train = x_train.reshape(x_train.shape[0], 28, 28, 1)
x_test = x_test.reshape(x_test.shape[0], 28, 28, 1)
#input_shape = (28, 28, 1)

# Making sure that the values are float so that we can get decimal points after
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')

# Normalizing the RGB codes by dividing it to the max RGB value.
x_train /= 255.0
x_test /= 255.0
print('x_train shape:', x_train.shape)
print('y_train shape:', y_train.shape)
print('x_test shape:', x_test.shape)
print('y_test shape:', y_test.shape)
```

x_train shape: (60000, 28, 28, 1)
y_train shape: (60000,)
x_test shape: (10000, 28, 28, 1)
y_test shape: (10000,)

```
In [8]: y_train[0:15]
```

```
Out[8]: array([5, 0, 4, 1, 9, 2, 1, 3, 1, 4, 3, 5, 3, 6, 1], dtype=uint8)
```

Definition of the neural network architecture

```
In [9]: # Creating a Sequential Model and adding the layers

def architecture(batch_normalization, dropout, input_shape, activation):
    model = Sequential()
    model.add(Conv2D(32, kernel_size=(5,5), input_shape=input_shape))

    model.add(MaxPooling2D(pool_size=(2, 2)))
    if batch_normalization:
        model.add(BatchNormalization())      #The recomendation is to perform ba

    model.add(Conv2D(64, kernel_size=(5,5), input_shape=input_shape))

    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Flatten()) # Flattening the 2D arrays for fully connected layer

    model.add(Dense(1024))
    if dropout:
        model.add(Dropout(0.2))
    if batch_normalization:
        model.add(BatchNormalization())  #The recomendation is to perform batc
    model.add(Activation(activation))

    model.add(Dense(10,activation='softmax'))

    return model
```

Generating a model of deep neural network

Playing with batch normalization and dropout, you will see that batch normalization improves better the network. Remember that batch normalization is applied before the activation.

Paper: Batch Normalization

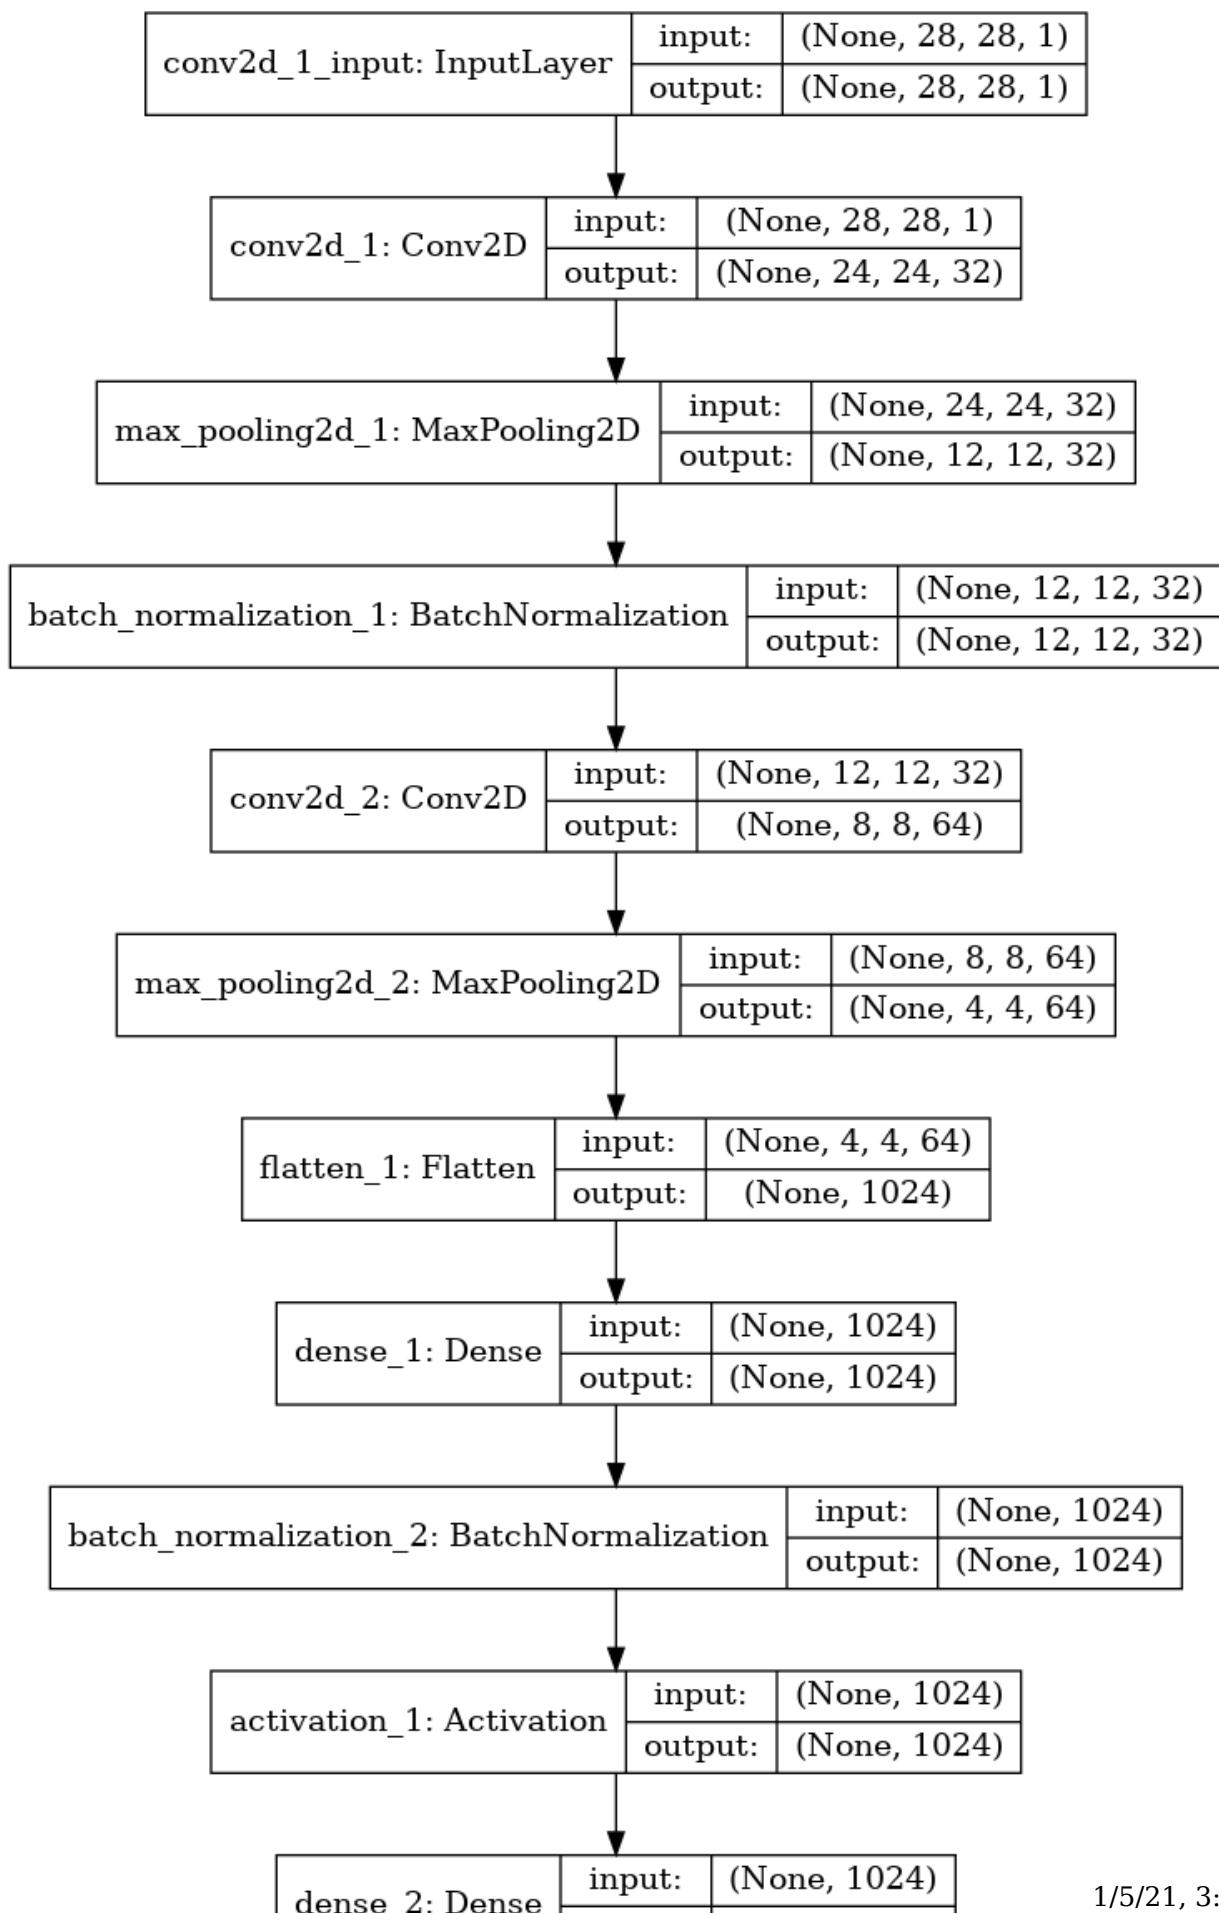
```
In [10]: batch_normalization=True
dropout=False
input_shape = (28, 28, 1)
activation = 'relu'

LeNet_model = architecture(batch_normalization, dropout, input_shape, activat
```

```
In [11]: # Plotting the architecture

plot_model(LeNet_model, to_file='LeNet.png', show_shapes=True, show_layer_nam
```

Out[11]:



In [12]: `LeNet_model.summary()`

Model: "sequential_1"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_1 (Conv2D)	(None, 24, 24, 32)	832
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 32)	0
batch_normalization_1 (Batch Normalization)	(None, 12, 12, 32)	128
conv2d_2 (Conv2D)	(None, 8, 8, 64)	51264
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 64)	0
flatten_1 (Flatten)	(None, 1024)	0
dense_1 (Dense)	(None, 1024)	1049600
batch_normalization_2 (Batch Normalization)	(None, 1024)	4096
activation_1 (Activation)	(None, 1024)	0
dense_2 (Dense)	(None, 10)	10250
<hr/>		
Total params: 1,116,170		
Trainable params: 1,114,058		
Non-trainable params: 2,112		

Keras: compiling methods

Compiling the model

In [13]: `#Compiling the model`

```
lr = 0.001
```

```
LeNet_model.compile(optimizer=optimizers.Adam(learning_rate=lr, beta_1=0.9, beta_2=0.999), loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

Running the model

In [14]: `start_time = time.time()`

```
num_epochs=20
```

```
history = LeNet_model.fit(x_train, y_train, batch_size=256, epochs=num_epochs)
```

```
end_time = time.time()
```

```
print("Time for training: {:.10f}s".format(end_time - start_time))
```

Train on 50400 samples, validate on 9600 samples

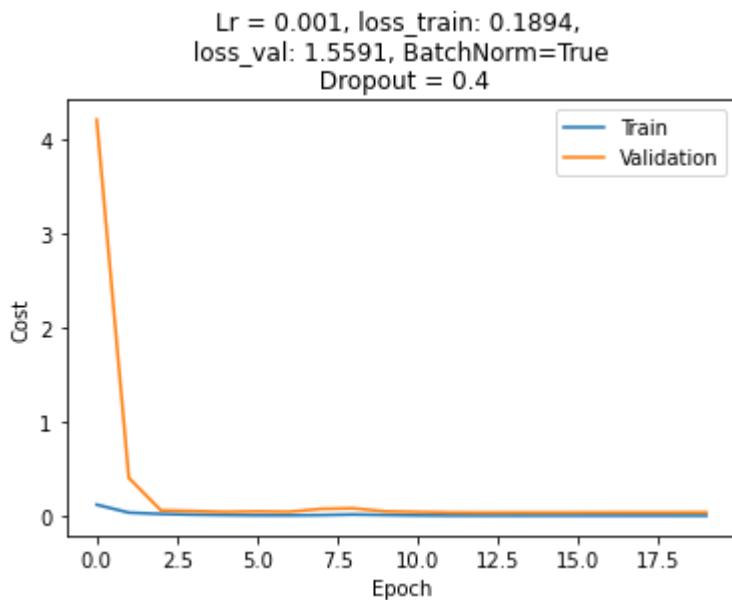
Epoch 1/20

1/5/21, 3:32 PM

```
50400/50400 [=====] - 3s 53us/step - loss: 0.1167 -  
accuracy: 0.9630 - val_loss: 4.2014 - val_accuracy: 0.1065  
Epoch 2/20  
50400/50400 [=====] - 1s 28us/step - loss: 0.0321 -  
accuracy: 0.9899 - val_loss: 0.4013 - val_accuracy: 0.8586  
Epoch 3/20  
50400/50400 [=====] - 1s 28us/step - loss: 0.0182 -  
accuracy: 0.9940 - val_loss: 0.0565 - val_accuracy: 0.9809  
Epoch 4/20  
50400/50400 [=====] - 1s 28us/step - loss: 0.0110 -  
accuracy: 0.9965 - val_loss: 0.0491 - val_accuracy: 0.9865  
Epoch 5/20  
50400/50400 [=====] - 1s 28us/step - loss: 0.0067 -  
accuracy: 0.9982 - val_loss: 0.0398 - val_accuracy: 0.9884  
Epoch 6/20  
50400/50400 [=====] - 1s 28us/step - loss: 0.0042 -  
accuracy: 0.9990 - val_loss: 0.0450 - val_accuracy: 0.9876  
Epoch 7/20  
50400/50400 [=====] - 1s 28us/step - loss: 0.0034 -  
accuracy: 0.9993 - val_loss: 0.0410 - val_accuracy: 0.9886  
Epoch 8/20  
50400/50400 [=====] - 1s 28us/step - loss: 0.0062 -  
accuracy: 0.9981 - val_loss: 0.0717 - val_accuracy: 0.9818  
Epoch 9/20  
50400/50400 [=====] - 1s 28us/step - loss: 0.0125 -  
accuracy: 0.9959 - val_loss: 0.0783 - val_accuracy: 0.9804  
Epoch 10/20  
50400/50400 [=====] - 1s 28us/step - loss: 0.0073 -  
accuracy: 0.9975 - val_loss: 0.0451 - val_accuracy: 0.9892  
Epoch 11/20  
50400/50400 [=====] - 1s 28us/step - loss: 0.0031 -  
accuracy: 0.9992 - val_loss: 0.0390 - val_accuracy: 0.9901  
Epoch 12/20  
50400/50400 [=====] - 1s 28us/step - loss: 7.1629e-0  
4 - accuracy: 0.9999 - val_loss: 0.0335 - val_accuracy: 0.9911  
Epoch 13/20  
50400/50400 [=====] - 1s 28us/step - loss: 2.4544e-0  
4 - accuracy: 1.0000 - val_loss: 0.0319 - val_accuracy: 0.9920  
Epoch 14/20  
50400/50400 [=====] - 1s 28us/step - loss: 1.7695e-0  
4 - accuracy: 1.0000 - val_loss: 0.0328 - val_accuracy: 0.9921  
Epoch 15/20  
50400/50400 [=====] - 1s 28us/step - loss: 1.2038e-0  
4 - accuracy: 1.0000 - val_loss: 0.0325 - val_accuracy: 0.9923  
Epoch 16/20  
50400/50400 [=====] - 1s 29us/step - loss: 1.0140e-0  
4 - accuracy: 1.0000 - val_loss: 0.0330 - val_accuracy: 0.9926  
Epoch 17/20  
50400/50400 [=====] - 1s 28us/step - loss: 9.0561e-0  
5 - accuracy: 1.0000 - val_loss: 0.0336 - val_accuracy: 0.9928  
Epoch 18/20  
50400/50400 [=====] - 1s 28us/step - loss: 7.2238e-0  
5 - accuracy: 1.0000 - val_loss: 0.0338 - val_accuracy: 0.9924  
Epoch 19/20  
50400/50400 [=====] - 1s 28us/step - loss: 6.3028e-0  
5 - accuracy: 1.0000 - val_loss: 0.0339 - val_accuracy: 0.9923  
Epoch 20/20  
50400/50400 [=====] - 1s 28us/step - loss: 5.3360e-0  
5 - accuracy: 1.0000 - val_loss: 0.0346 - val_accuracy: 0.9922
```

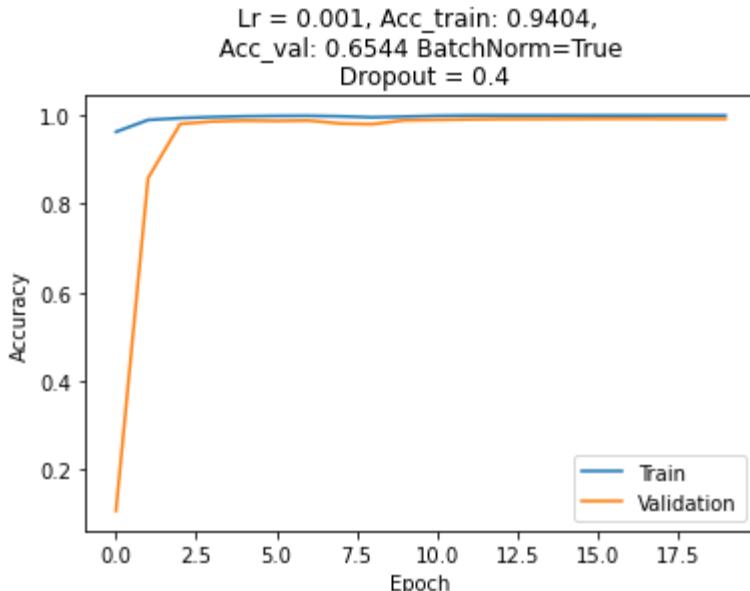
Plotting the loss function

```
In [15]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Lr = 0.001, loss_train: 0.1894, \n loss_val: 1.5591, BatchNorm=True')
plt.ylabel('Cost')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper right')
# plt.ylim(top=13)
# plt.ylim(bottom=0)
plt.show()
```



Plotting the accuracy

```
In [16]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Lr = 0.001, Acc_train: 0.9404, \n Acc_val: 0.6544 BatchNorm=True \n')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='lower right')
plt.show()
```



```
In [17]: # Predicting the image associated to the each sample in the test set (X_test)
predictions = LeNet_model.predict(x_test)
```

```
In [18]: print(type(predictions))
print(predictions.shape)
```

```
<class 'numpy.ndarray'>
(10000, 10)
```

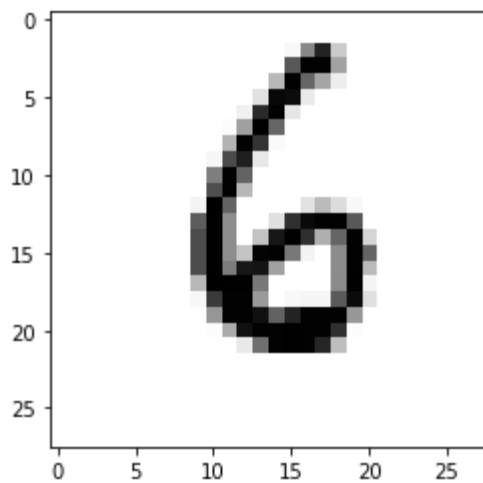
```
In [19]: sample = 91
print(predictions[sample])
print("\nPredicted digit:", np.argmax(predictions[sample]))
```

```
[5.7965465e-13 1.8259214e-12 4.5605764e-10 4.8307799e-12 3.7656635e-11
 6.1520482e-09 1.0000000e+00 1.1678162e-16 1.6059101e-10 9.9855061e-14]
```

Predicted digit: 6

Displaying the image associated to this sample.

```
In [20]: plt.imshow(x_test[sample], cmap='Greys')
plt.show()
```



Deep Learning: AlexNet

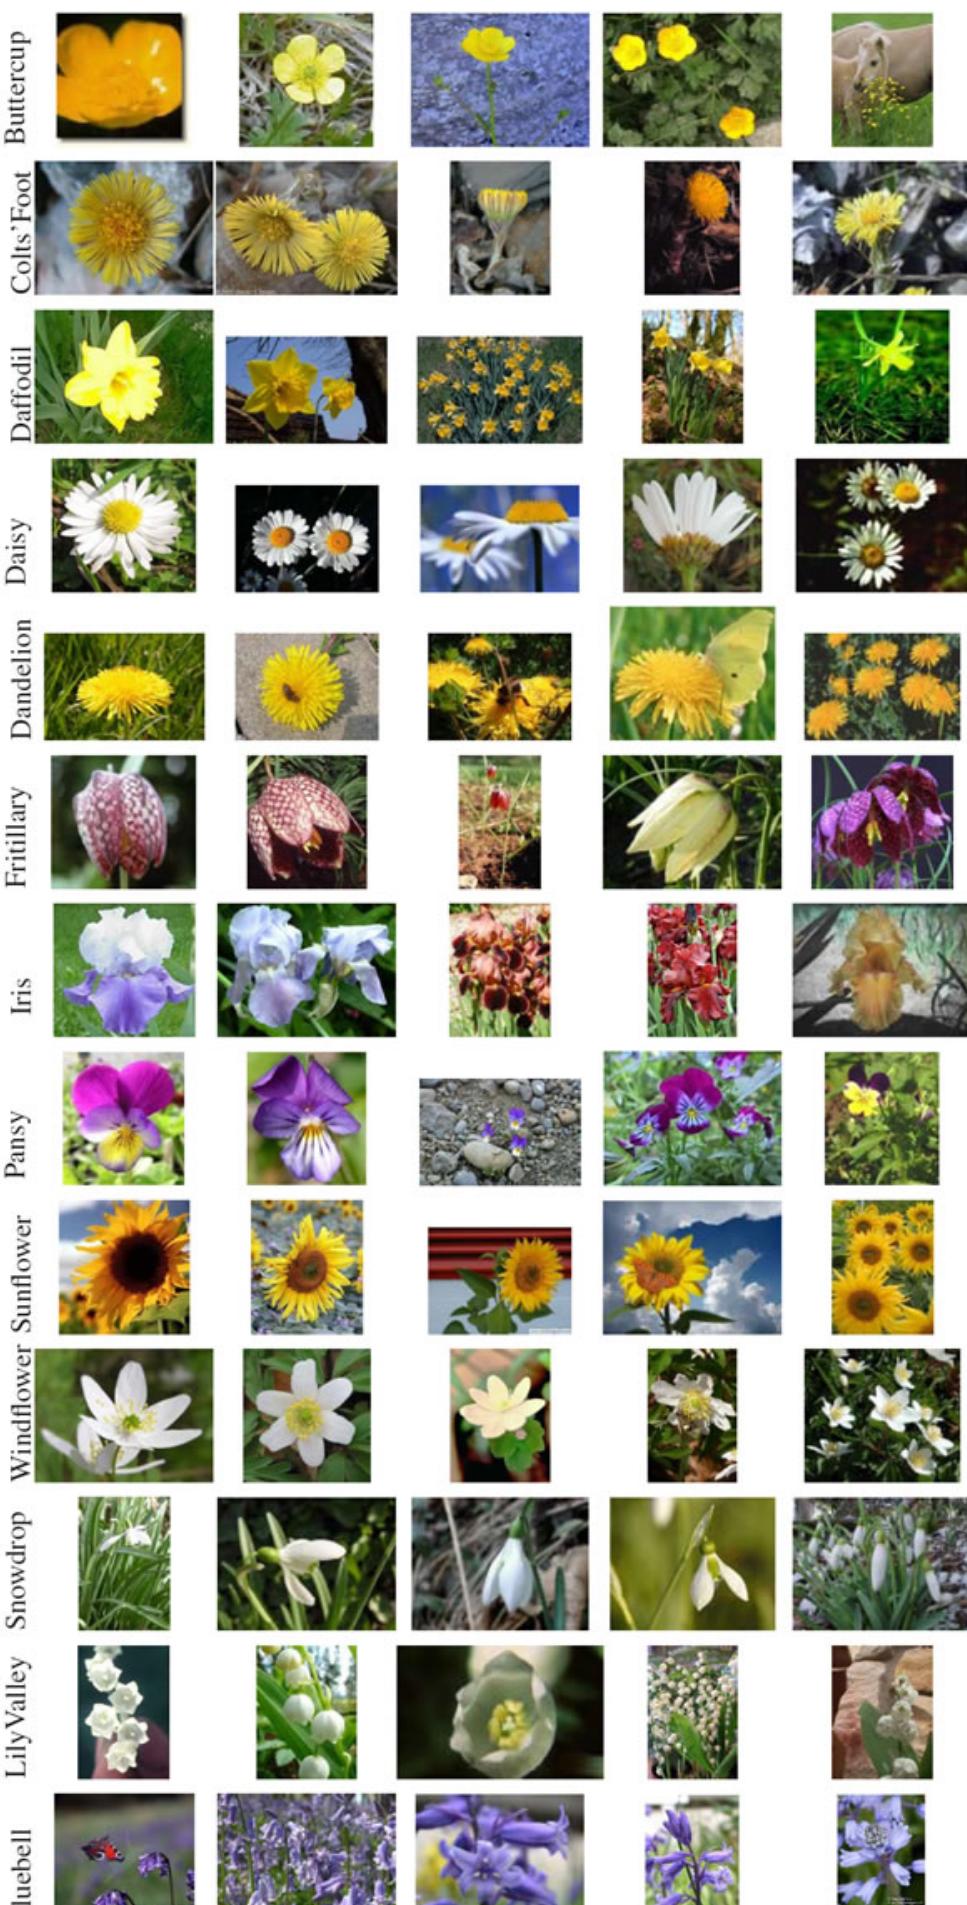
```
In [21]: import numpy as np
import matplotlib.pyplot as plt
import time

from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Activation, Dense, Flatten
from keras.layers import Activation, Dropout, BatchNormalization
from keras.utils import plot_model
from keras import optimizers

np.random.seed(10)
```

Data of the System to be analyzed: oxford17

The [oxford17 database](#)



Generation or extraction of the raw data

Install the library tflearn to get the data.

TFLearn library

```
In [22]: import tensorflow as tf
tf.set_random_seed(42)
import tflearn
from tflearn.data_utils import image_preloader
```

```
WARNING:tensorflow:From /home/bokhimi/anaconda3/envs/tf-gpu/lib/python3.8/site-packages/tensorflow/python/compat/v2_compat.py:96: disable_resource_variables (from tensorflow.python.ops.variable_scope) is deprecated and will be removed in a future version.
```

Instructions for updating:

non-resource variables are not supported in the long term

Analysis of the raw data

The oxford17 dataset consists of 1360 colour images (224 pixels high and 224 pixels width) of flowers in 17 classes, with 80 images per class. All images will be used for training. Before running the model, it will be indicated the ratio of samples that will be used for validation.

The 17 classes are:

index	class name
0	Daffodil
1	Snowdrop
2	Daisy
3	ColtsFoot
4	Dandelion
5	Cowslip
6	Buttercup
7	Windflower
8	Pansy
9	LilyValley
10	Bluebell
11	Crocus
12	Iris

index	class name
13	Tigerlily
14	Tulip
15	Fritillary

Viewing one sample from the data sets

We define a dictionary to associate the class number to a class name.

```
In [23]: dic = {0: 'Daffodil', 1: 'Snowdrop', 2: 'Daisy', 3: 'ColtsFoot', 4: 'Dandelio',
5: 'Cowslip', 6: 'Buttercup', 7: 'Windflower', 8: 'Pansy', 9: 'LilyVall',
10: 'Bluebell', 11: 'Crocus', 12: 'Iris', 13: 'Tigerlily', 14: 'Tulip',
15: 'Fritillary', 16: 'Sunflower'}
```

Next, we show a sample: its target and image.

```
In [24]: # Plotting the content of a sample

sample = 72

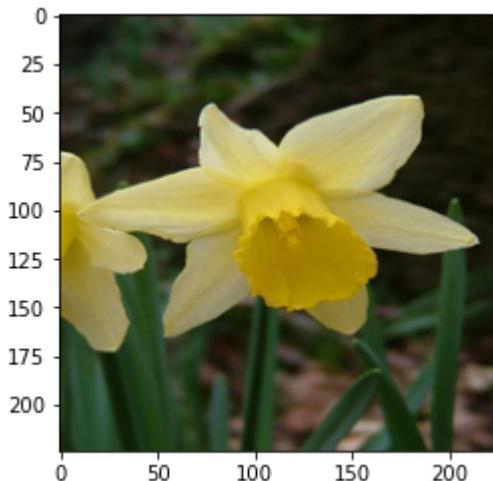
plt.imshow(train_x[sample]);
print('y =', np.squeeze(train_y[sample]))

for i in [i for i,x in enumerate(train_y[sample]) if x == 1]:
    print('')

print('y =', i, ';', 'the sample', sample, 'corresponds to a(an)', dic[i])

y = [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
```

y = 0 ; the sample 72 corresponds to a(an) Daffodil



Transformation of the raw data

```
In [25]: print('the shape is', train_x.shape)  
the shape is (1360, 224, 224, 3)
```

```
In [26]: print(train_x[0][0:5][0:2])  
[[[0.06666667 0.05882353 0.10980392]  
[0.09019608 0.08235294 0.13333334]  
[0.13333334 0.1254902 0.1764706 ]  
...  
[0.31764707 0.31764707 0.31764707]  
[0.30980393 0.30980393 0.30980393]  
[0.29411766 0.29411766 0.29411766]]  
  
[[0.12941177 0.12156863 0.17254902]  
[0.10588235 0.09803922 0.14901961]  
[0.09803922 0.09019608 0.14117648]  
...  
[0.30980393 0.30980393 0.30980393]  
[0.23529412 0.23529412 0.23529412]  
[0.24313726 0.24313726 0.24313726]]]
```

```
In [27]: print(train_y[0])  
[0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
```

Transformation of the raw data

The raw data are jet renormalized. We do not do anything more

```
In [28]: print('train_x shape:', train_x.shape)  
print('train_y shape:', train_y.shape)  
  
train_x shape: (1360, 224, 224, 3)  
train_y shape: (1360, 17)
```

```
In [29]: from sklearn.model_selection import train_test_split  
  
# Choose your test size to split between training and testing sets:  
train_x, test_x, train_y, test_y = train_test_split(train_x, train_y, test_size=0.2)
```

```
In [30]: print(train_x.shape)  
print(test_x.shape)  
print(train_y.shape)  
print(test_y.shape)  
  
(1224, 224, 224, 3)  
(136, 224, 224, 3)  
(1224, 17)  
(136, 17)
```

Definition of the neural network architecture

Keras has two different modes to define the architecture:

1. The sequential model. It is a sequential stack of layers.
2. The functional API. It is the way to go for defining complex models, such as multi-output models, directed acyclic graphs, or models with shared layers.

In the present case, we will use the sequential mode for constructing the architecture of the network.

[Keras: Sequential model API](#)

[Keras: Convolutional layers](#)

[Keras: Pooling layers](#)

[Keras: Batch Normalization](#)

```
In [31]: # Creating a Sequential Model and adding the layers

def architecture(batch_normalization, dropout, input_shape, activation):

    # Creating a sequential model
    model = Sequential()

    # 1st Convolutional layer
    model.add(Conv2D(filters=96, activation=activation, input_shape=input_shape,
                     kernel_size=(11,11), strides=(4,4), padding='valid', kernel_initializer='he_uniform'))
    # Pooling
    model.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='valid'))
    if batch_normalization:
        model.add(BatchNormalization())

    # 2nd Convolutional Layer
    model.add(Conv2D(filters=256, activation=activation, kernel_size=(5,5),
                     strides=(1,1), padding='valid'))
    # Pooling
    model.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='valid'))
    if batch_normalization:
        model.add(BatchNormalization())

    # 3rd Convolutional Layer
    model.add(Conv2D(filters=384, activation=activation, kernel_size=(3,3),
                     strides=(1,1), padding='valid'))

    # 4th Convolutional Layer
    model.add(Conv2D(filters=384, activation=activation, kernel_size=(3,3),
                     strides=(1,1), padding='valid'))

    # 5th Convolutional Layer
    model.add(Conv2D(filters=256, activation=activation, kernel_size=(3,3),
                     strides=(1,1), padding='valid'))
    # Pooling
    model.add(MaxPooling2D(pool_size=(2,2), strides=(2,2), padding='valid'))
    if batch_normalization:
        model.add(BatchNormalization())

    # Passing it to a dense layer
    model.add(Flatten())
    if dropout:
        model.add(Dropout(0.4))

    # 1st Dense Layer
    model.add(Dense(512, activation=activation, input_shape=(224*224*3,), kernel_initializer='he_uniform'))
    # Add Dropout to prevent overfitting
    if dropout:
        model.add(Dropout(0.4))
    if batch_normalization:
        model.add(BatchNormalization())

    # 2nd Dense Layer
    model.add(Dense(512, activation=activation, kernel_initializer = 'he_uniform'))
    model.add(Activation('relu'))
    # Add Dropout
    if dropout:
        model.add(Dropout(0.4))
    if batch_normalization:
        model.add(BatchNormalization())
```

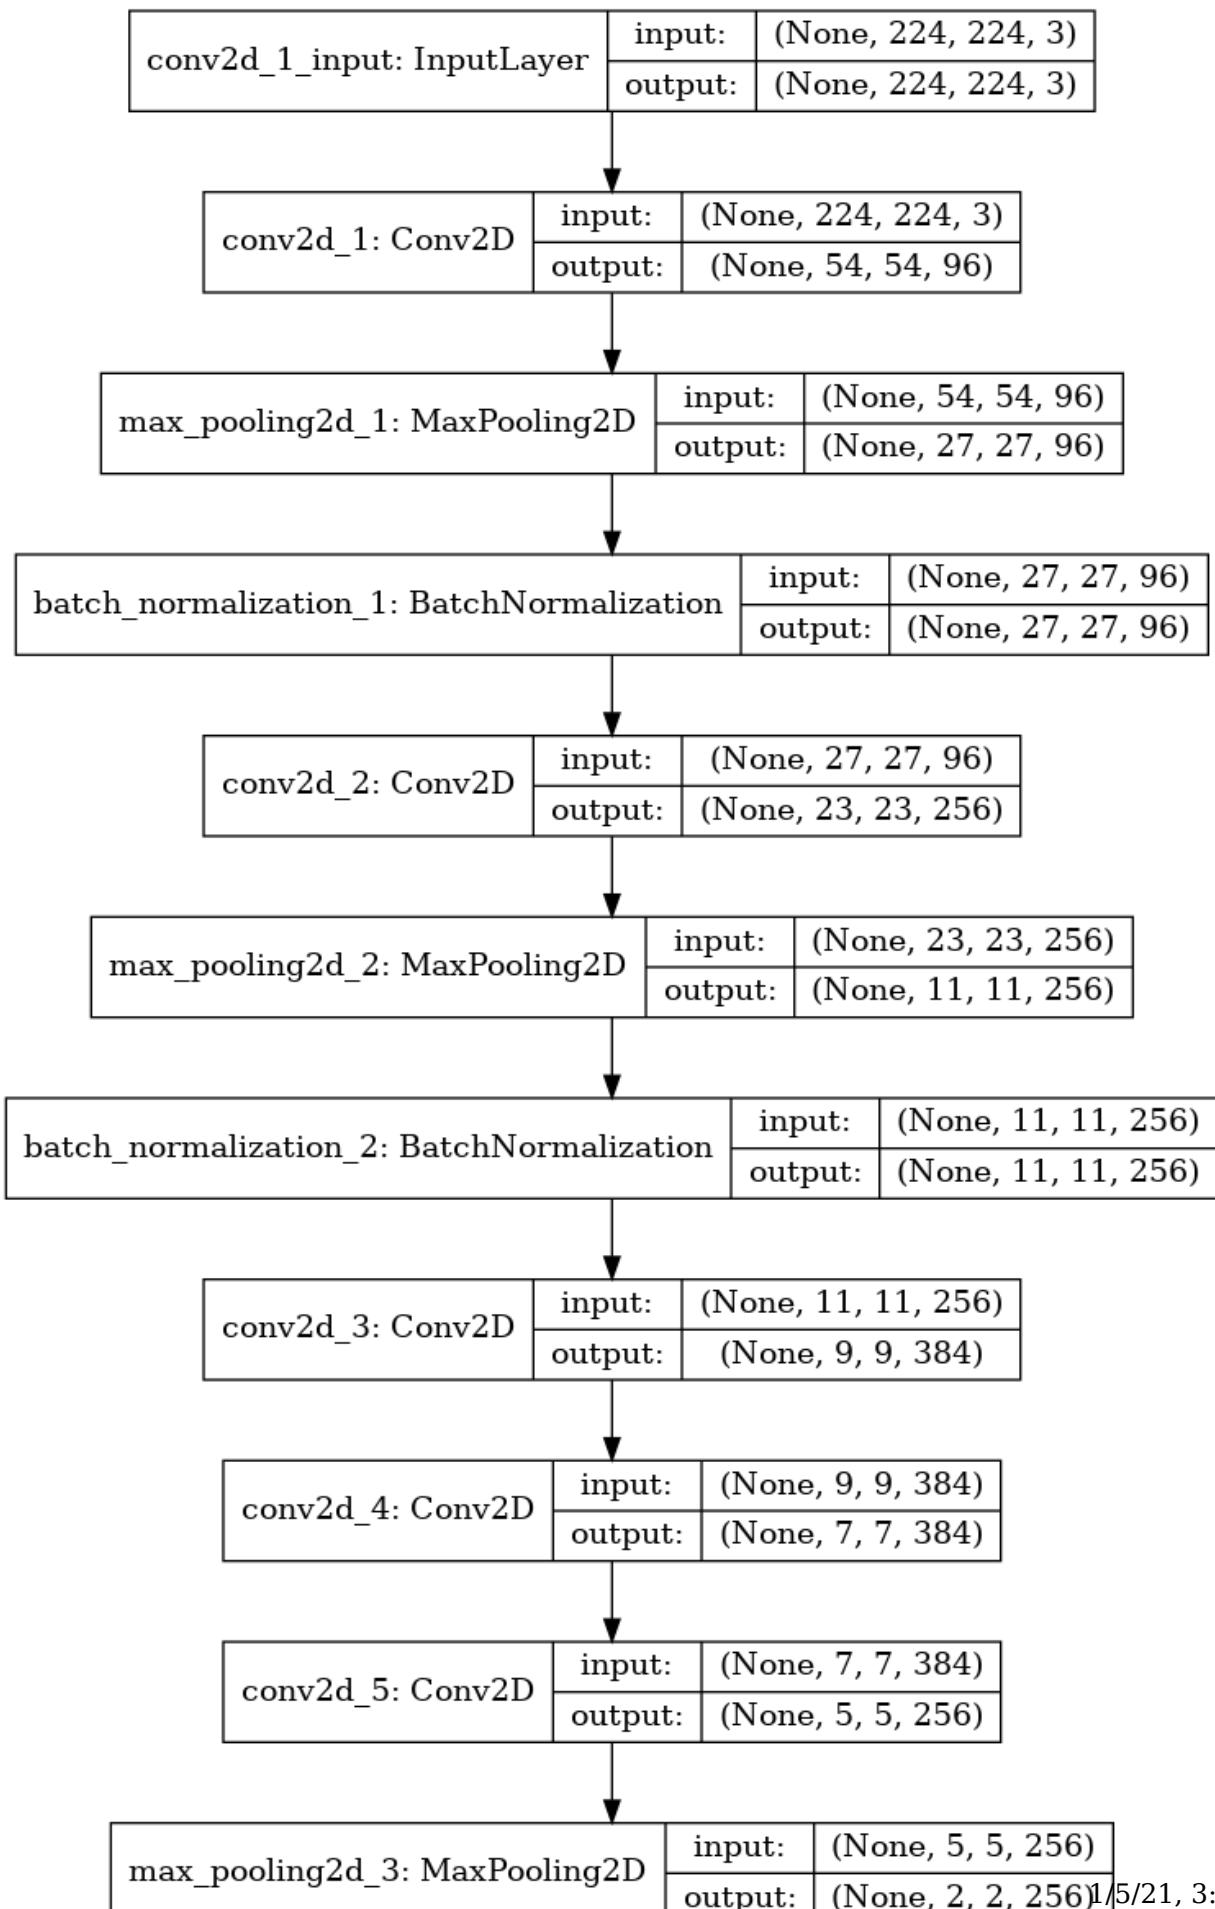
Generating a model of deep neural network

```
In [32]: # Generating the model using the defined architecture  
  
batch_normalization=True  
dropout=True  
one_image = (224, 224, 3)  
activation = 'relu'  
  
oxflower17_model = architecture(batch_normalization, dropout, one_image, acti
```

WARNING:tensorflow:From /home/bokhimi/anaconda3/envs/tf-gpu/lib/python3.8/site-packages/tensorflow/python/ops/resource_variable_ops.py:1659: calling BaseResourceVariable.__init__ (from tensorflow.python.ops.resource_variable_ops) with constraint is deprecated and will be removed in a future version.
Instructions for updating:
If using Keras pass *_constraint arguments to layers.

```
In [33]: plot_model(oxflower17_model, to_file='oxflower17_model.png', show_shapes=True)
```

Out[33]:



```
In [34]: oxflower17_model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 54, 54, 96)	34944
max_pooling2d_1 (MaxPooling2D)	(None, 27, 27, 96)	0
batch_normalization_1 (BatchNormalization)	(None, 27, 27, 96)	384
conv2d_2 (Conv2D)	(None, 23, 23, 256)	614656
max_pooling2d_2 (MaxPooling2D)	(None, 11, 11, 256)	0
batch_normalization_2 (BatchNormalization)	(None, 11, 11, 256)	1024
conv2d_3 (Conv2D)	(None, 9, 9, 384)	885120
conv2d_4 (Conv2D)	(None, 7, 7, 384)	1327488
conv2d_5 (Conv2D)	(None, 5, 5, 256)	884992
max_pooling2d_3 (MaxPooling2D)	(None, 2, 2, 256)	0
batch_normalization_3 (BatchNormalization)	(None, 2, 2, 256)	1024
flatten_1 (Flatten)	(None, 1024)	0
dropout_1 (Dropout)	(None, 1024)	0
dense_1 (Dense)	(None, 512)	524800
dropout_2 (Dropout)	(None, 512)	0
batch_normalization_4 (BatchNormalization)	(None, 512)	2048
dense_2 (Dense)	(None, 512)	262656
activation_1 (Activation)	(None, 512)	0
dropout_3 (Dropout)	(None, 512)	0
batch_normalization_5 (BatchNormalization)	(None, 512)	2048
dense_3 (Dense)	(None, 17)	8721
<hr/>		
Total params: 4,549,905		
Trainable params: 4,546,641		
Non-trainable params: 3,264		

Compiling the model

```
In [35]: #Compiling the model using Adam as optimizer

lr = 0.001 # Learning rate

oxflower17_model.compile(loss='categorical_crossentropy', metrics=['accuracy']
optimizer=optimizers.Adam(learning_rate=lr, beta_1=0.9, beta_2=0.999, amsgrad=
```

Running the model

```
In [36]: start_time = time.time()

batch_size=32
num_epochs = 50

history = oxfower17_model.fit(train_x, train_y, batch_size=batch_size,
epochs=num_epochs, validation_data=(test_x,test_y),verbose=1, shuffle=1)

end_time = time.time()
print("Time for training: {:.10f}s".format(end_time - start_time))
```

Train on 1224 samples, validate on 136 samples
Epoch 1/50
1224/1224 [=====] - 2s 2ms/step - loss: 2.9391 - accuracy: 0.1462 - val_loss: 40.2617 - val_accuracy: 0.0441
Epoch 2/50
1224/1224 [=====] - 1s 1ms/step - loss: 2.4698 - accuracy: 0.2451 - val_loss: 5.7297 - val_accuracy: 0.1471
Epoch 3/50
1224/1224 [=====] - 1s 1ms/step - loss: 2.0694 - accuracy: 0.3619 - val_loss: 4.0435 - val_accuracy: 0.2647
Epoch 4/50
1224/1224 [=====] - 1s 1ms/step - loss: 1.9486 - accuracy: 0.3660 - val_loss: 2.3372 - val_accuracy: 0.2868
Epoch 5/50
1224/1224 [=====] - 1s 1ms/step - loss: 1.8283 - accuracy: 0.4044 - val_loss: 2.5977 - val_accuracy: 0.3015
Epoch 6/50
1224/1224 [=====] - 1s 1ms/step - loss: 1.7595 - accuracy: 0.4257 - val_loss: 1.9638 - val_accuracy: 0.4559
Epoch 7/50
1224/1224 [=====] - 1s 1ms/step - loss: 1.6529 - accuracy: 0.4371 - val_loss: 1.9267 - val_accuracy: 0.4191
Epoch 8/50
1224/1224 [=====] - 1s 1ms/step - loss: 1.5097 - accuracy: 0.4959 - val_loss: 2.0047 - val_accuracy: 0.4706
Epoch 9/50
1224/1224 [=====] - 1s 1ms/step - loss: 1.4617 - accuracy: 0.5016 - val_loss: 2.3358 - val_accuracy: 0.3235
Epoch 10/50
1224/1224 [=====] - 1s 1ms/step - loss: 1.4467 - accuracy: 0.5049 - val_loss: 2.1216 - val_accuracy: 0.3676
Epoch 11/50
1224/1224 [=====] - 1s 1ms/step - loss: 1.4191 - accuracy: 0.5204 - val_loss: 1.6400 - val_accuracy: 0.4779
Epoch 12/50
1224/1224 [=====] - 1s 1ms/step - loss: 1.3123 - accuracy: 0.5204 - val_loss: 1.6400 - val_accuracy: 0.4779

```
uracy: 0.5556 - val_loss: 5.1617 - val_accuracy: 0.1691
Epoch 13/50
1224/1224 [=====] - 1s 1ms/step - loss: 1.3954 - acc
uracy: 0.5172 - val_loss: 2.4555 - val_accuracy: 0.3529
Epoch 14/50
1224/1224 [=====] - 1s 1ms/step - loss: 1.4642 - acc
uracy: 0.5131 - val_loss: 2.7454 - val_accuracy: 0.3235
Epoch 15/50
1224/1224 [=====] - 1s 1ms/step - loss: 1.2348 - acc
uracy: 0.5662 - val_loss: 1.5613 - val_accuracy: 0.5662
Epoch 16/50
1224/1224 [=====] - 1s 1ms/step - loss: 1.2445 - acc
uracy: 0.5801 - val_loss: 2.3547 - val_accuracy: 0.3971
Epoch 17/50
1224/1224 [=====] - 1s 1ms/step - loss: 1.0969 - acc
uracy: 0.6291 - val_loss: 1.5511 - val_accuracy: 0.4559
Epoch 18/50
1224/1224 [=====] - 1s 1ms/step - loss: 1.0636 - acc
uracy: 0.6389 - val_loss: 2.2117 - val_accuracy: 0.4706
Epoch 19/50
1224/1224 [=====] - 1s 1ms/step - loss: 1.0597 - acc
uracy: 0.6291 - val_loss: 1.4577 - val_accuracy: 0.5221
Epoch 20/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.9872 - acc
uracy: 0.6560 - val_loss: 1.6981 - val_accuracy: 0.5147
Epoch 21/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.9452 - acc
uracy: 0.6871 - val_loss: 2.0846 - val_accuracy: 0.4118
Epoch 22/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.8678 - acc
uracy: 0.7059 - val_loss: 1.7226 - val_accuracy: 0.4926
Epoch 23/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.7931 - acc
uracy: 0.7181 - val_loss: 2.0558 - val_accuracy: 0.4559
Epoch 24/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.7258 - acc
uracy: 0.7443 - val_loss: 1.5422 - val_accuracy: 0.5368
Epoch 25/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.7646 - acc
uracy: 0.7279 - val_loss: 1.4043 - val_accuracy: 0.5588
Epoch 26/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.6794 - acc
uracy: 0.7508 - val_loss: 2.1713 - val_accuracy: 0.4265
Epoch 27/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.8126 - acc
uracy: 0.7288 - val_loss: 1.9466 - val_accuracy: 0.4265
Epoch 28/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.5905 - acc
uracy: 0.7917 - val_loss: 1.4886 - val_accuracy: 0.5956
Epoch 29/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.5326 - acc
uracy: 0.8276 - val_loss: 1.8373 - val_accuracy: 0.5368
Epoch 30/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.7779 - acc
uracy: 0.7484 - val_loss: 1.8009 - val_accuracy: 0.4706
Epoch 31/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.6135 - acc
uracy: 0.7884 - val_loss: 1.5818 - val_accuracy: 0.5368
Epoch 32/50
```

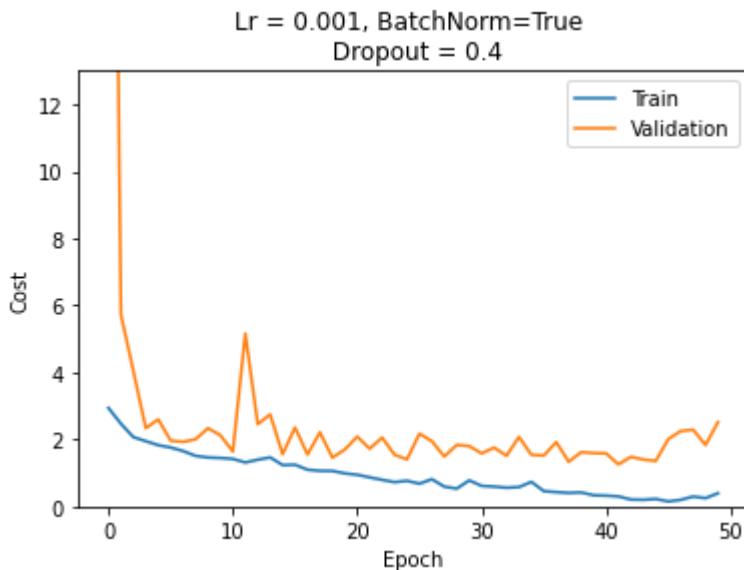
```
1224/1224 [=====] - 1s 1ms/step - loss: 0.5922 - accuracy: 0.7917 - val_loss: 1.7567 - val_accuracy: 0.5000
Epoch 33/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.5589 - accuracy: 0.8047 - val_loss: 1.5121 - val_accuracy: 0.5662
Epoch 34/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.5783 - accuracy: 0.7998 - val_loss: 2.0773 - val_accuracy: 0.5662
Epoch 35/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.7335 - accuracy: 0.7557 - val_loss: 1.5412 - val_accuracy: 0.5882
Epoch 36/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.4572 - accuracy: 0.8480 - val_loss: 1.5193 - val_accuracy: 0.5735
Epoch 37/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.4279 - accuracy: 0.8578 - val_loss: 1.9197 - val_accuracy: 0.5588
Epoch 38/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.4046 - accuracy: 0.8529 - val_loss: 1.3339 - val_accuracy: 0.6176
Epoch 39/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.4177 - accuracy: 0.8546 - val_loss: 1.6187 - val_accuracy: 0.5588
Epoch 40/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.3340 - accuracy: 0.8946 - val_loss: 1.5935 - val_accuracy: 0.5074
Epoch 41/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.3233 - accuracy: 0.8815 - val_loss: 1.5850 - val_accuracy: 0.6103
Epoch 42/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.2993 - accuracy: 0.9101 - val_loss: 1.2564 - val_accuracy: 0.6103
Epoch 43/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.2061 - accuracy: 0.9346 - val_loss: 1.4779 - val_accuracy: 0.6103
Epoch 44/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.1994 - accuracy: 0.9314 - val_loss: 1.4005 - val_accuracy: 0.7132
Epoch 45/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.2234 - accuracy: 0.9289 - val_loss: 1.3579 - val_accuracy: 0.6324
Epoch 46/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.1498 - accuracy: 0.9518 - val_loss: 2.0061 - val_accuracy: 0.5515
Epoch 47/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.1963 - accuracy: 0.9322 - val_loss: 2.2492 - val_accuracy: 0.5368
Epoch 48/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.2955 - accuracy: 0.9118 - val_loss: 2.2917 - val_accuracy: 0.4926
Epoch 49/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.2477 - accuracy: 0.9232 - val_loss: 1.8336 - val_accuracy: 0.6618
Epoch 50/50
1224/1224 [=====] - 1s 1ms/step - loss: 0.3940 - accuracy: 0.8701 - val_loss: 2.5182 - val_accuracy: 0.5117
```

- Note: if you run `fit()` again, the model will continue training,

starting with the parameters it has already learnt, instead of reinitializing them.

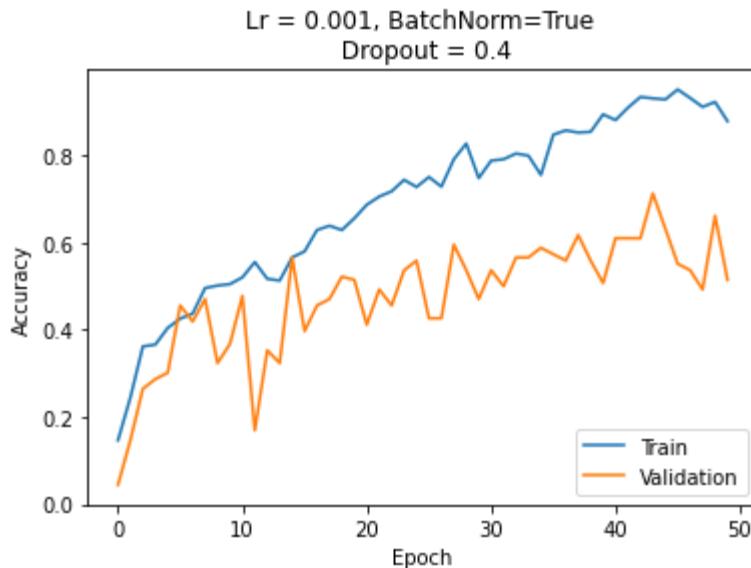
Plotting the loss function

```
In [37]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Lr = 0.001, BatchNorm=True \n Dropout = 0.4')
plt.ylabel('Cost')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper right')
plt.ylim(top=13)      # The instruction is used to limit the upper value of the
plt.ylim(bottom=0)    # The instruction is used to limit the lower value of the
plt.show()
```



Plotting the accuracy

```
In [38]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Lr = 0.001, BatchNorm=True \n Dropout = 0.4')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='lower right')
plt.show()
```



Data augmentation

`shear_range`, `zoom_range`, and `horizontal_flip` are some of the parameters available in Keras that define the transformation of the images

[Keras: Data augmentation](#)

Generating a model of deep neural network

```
In [39]: # Generating the model using the defined architecture  
batch_normalization=True  
dropout=True  
one_image = (224, 224, 3)  
activation = 'relu'  
  
oxflower17_model = architecture(batch_normalization, dropout, one_image, acti
```

Compiling the model

```
In [40]: #Compiling the model using Adam as optimizer  
  
lr = 0.001 # Learning rate  
  
oxflower17_model.compile(loss='categorical_crossentropy', metrics=['accuracy'])  
optimizer=optimizers.Adam(learning_rate=lr, beta_1=0.9, beta_2=0.999, amsgrad=
```

```
In [41]: from keras.preprocessing.image import ImageDataGenerator

train_datagen = ImageDataGenerator(
    shear_range = 0.2,
    zoom_range = 0.2,
    horizontal_flip = True
)
```

Running the model

Comment: Keras flow method

This process requires long times, depending of the number of steps per epoch, the number of epochs and the number of images that will be generated during the data augmentation (batch_size)

```
In [42]: # compute quantities required for featurewise normalization
# (std, mean, and principal components if ZCA whitening is applied)
train_datagen.fit(train_x)

train_generator = train_datagen.flow(
    train_x,
    train_y,
    batch_size = 32,
    shuffle=True
)
```

```
In [43]: get_steps_augment = 64

print ("X_train shape: " + str(train_x.shape[0]))
steps = int(train_x.shape[0]/get_steps_augment)
print("Augmentation steps = {}".format(steps))
```

X_train shape: 1224
 Augmentation steps = 19

```
In [44]: start_time = time.time()

num_epochs = 100

history = oxford17_model.fit(train_generator, steps_per_epoch=steps,\n                           epochs=num_epochs, validation_data=(test_x, test_y), verbose=1, shuff

end_time = time.time()
print("Time for training: {:.4f}s".format(end_time - start_time))

Epoch 1/100
19/19 [=====] - 3s 176ms/step - loss: 3.1353 - accuracy: 0.1168 - val_loss: 44.1296 - val_accuracy: 0.0515
Epoch 2/100
19/19 [=====] - 3s 184ms/step - loss: 2.6787 - accuracy: 0.1764 - val_loss: 23.8188 - val_accuracy: 0.1029
```

```
Epoch 3/100
19/19 [=====] - 3s 177ms/step - loss: 2.5314 - accuracy: 0.2346 - val_loss: 8.1675 - val_accuracy: 0.1471
Epoch 4/100
19/19 [=====] - 3s 182ms/step - loss: 2.3450 - accuracy: 0.2681 - val_loss: 6.4789 - val_accuracy: 0.1324
Epoch 5/100
19/19 [=====] - 4s 188ms/step - loss: 2.1867 - accuracy: 0.3240 - val_loss: 4.2634 - val_accuracy: 0.1691
Epoch 6/100
19/19 [=====] - 3s 179ms/step - loss: 2.0995 - accuracy: 0.3373 - val_loss: 2.7618 - val_accuracy: 0.2426
Epoch 7/100
19/19 [=====] - 4s 188ms/step - loss: 2.1194 - accuracy: 0.3141 - val_loss: 2.4809 - val_accuracy: 0.2868
Epoch 8/100
19/19 [=====] - 3s 178ms/step - loss: 2.1393 - accuracy: 0.3339 - val_loss: 2.3831 - val_accuracy: 0.3162
Epoch 9/100
19/19 [=====] - 3s 176ms/step - loss: 1.9776 - accuracy: 0.3613 - val_loss: 1.7023 - val_accuracy: 0.4779
Epoch 10/100
19/19 [=====] - 3s 177ms/step - loss: 1.8667 - accuracy: 0.3832 - val_loss: 1.9932 - val_accuracy: 0.3676
Epoch 11/100
19/19 [=====] - 4s 190ms/step - loss: 1.8851 - accuracy: 0.3832 - val_loss: 1.9633 - val_accuracy: 0.3897
Epoch 12/100
19/19 [=====] - 3s 173ms/step - loss: 1.7090 - accuracy: 0.4195 - val_loss: 2.2833 - val_accuracy: 0.3015
Epoch 13/100
19/19 [=====] - 3s 181ms/step - loss: 1.8009 - accuracy: 0.3921 - val_loss: 2.4853 - val_accuracy: 0.2647
Epoch 14/100
19/19 [=====] - 4s 184ms/step - loss: 1.7312 - accuracy: 0.4095 - val_loss: 2.4010 - val_accuracy: 0.3456
Epoch 15/100
19/19 [=====] - 3s 182ms/step - loss: 1.6900 - accuracy: 0.4161 - val_loss: 2.1207 - val_accuracy: 0.3676
Epoch 16/100
19/19 [=====] - 3s 178ms/step - loss: 1.6009 - accuracy: 0.4391 - val_loss: 1.6227 - val_accuracy: 0.4412
Epoch 17/100
19/19 [=====] - 3s 167ms/step - loss: 1.8838 - accuracy: 0.4304 - val_loss: 2.3893 - val_accuracy: 0.3750
Epoch 18/100
19/19 [=====] - 3s 184ms/step - loss: 1.8702 - accuracy: 0.3799 - val_loss: 2.2472 - val_accuracy: 0.3603
Epoch 19/100
19/19 [=====] - 4s 204ms/step - loss: 1.5802 - accuracy: 0.4523 - val_loss: 2.1612 - val_accuracy: 0.3971
Epoch 20/100
19/19 [=====] - 3s 179ms/step - loss: 1.6811 - accuracy: 0.4401 - val_loss: 1.8556 - val_accuracy: 0.4706
Epoch 21/100
19/19 [=====] - 3s 180ms/step - loss: 1.6443 - accuracy: 0.4293 - val_loss: 1.9464 - val_accuracy: 0.4632
Epoch 22/100
19/19 [=====] - 3s 176ms/step - loss: 1.5040 - accuracy: 0.4500 - val_loss: 1.7856 - val_accuracy: 0.4824
```

```
acy: 0.4984 - val_loss: 1.8450 - val_accuracy: 0.4706
Epoch 23/100
19/19 [=====] - 3s 173ms/step - loss: 1.5123 - accuracy: 0.4880 - val_loss: 2.0540 - val_accuracy: 0.3382
Epoch 24/100
19/19 [=====] - 3s 177ms/step - loss: 1.5558 - accuracy: 0.4675 - val_loss: 2.4349 - val_accuracy: 0.3603
Epoch 25/100
19/19 [=====] - 3s 179ms/step - loss: 1.5712 - accuracy: 0.4589 - val_loss: 2.2818 - val_accuracy: 0.4265
Epoch 26/100
19/19 [=====] - 4s 202ms/step - loss: 1.4107 - accuracy: 0.5082 - val_loss: 2.5848 - val_accuracy: 0.3824
Epoch 27/100
19/19 [=====] - 3s 182ms/step - loss: 1.4136 - accuracy: 0.5214 - val_loss: 1.5447 - val_accuracy: 0.5221
Epoch 28/100
19/19 [=====] - 3s 169ms/step - loss: 1.5443 - accuracy: 0.4897 - val_loss: 2.1851 - val_accuracy: 0.3750
Epoch 29/100
19/19 [=====] - 3s 177ms/step - loss: 1.3854 - accuracy: 0.5378 - val_loss: 1.7810 - val_accuracy: 0.4338
Epoch 30/100
19/19 [=====] - 3s 175ms/step - loss: 1.3356 - accuracy: 0.5493 - val_loss: 1.8170 - val_accuracy: 0.4485
Epoch 31/100
19/19 [=====] - 3s 171ms/step - loss: 1.3850 - accuracy: 0.5154 - val_loss: 1.4959 - val_accuracy: 0.5368
Epoch 32/100
19/19 [=====] - 3s 176ms/step - loss: 1.3862 - accuracy: 0.5137 - val_loss: 4.7515 - val_accuracy: 0.2279
Epoch 33/100
19/19 [=====] - 4s 186ms/step - loss: 1.3504 - accuracy: 0.5296 - val_loss: 2.4606 - val_accuracy: 0.3897
Epoch 34/100
19/19 [=====] - 3s 169ms/step - loss: 1.3227 - accuracy: 0.5565 - val_loss: 1.7868 - val_accuracy: 0.4559
Epoch 35/100
19/19 [=====] - 3s 181ms/step - loss: 1.3556 - accuracy: 0.5428 - val_loss: 2.0456 - val_accuracy: 0.3750
Epoch 36/100
19/19 [=====] - 3s 177ms/step - loss: 1.2548 - accuracy: 0.5493 - val_loss: 2.9642 - val_accuracy: 0.3015
Epoch 37/100
19/19 [=====] - 3s 171ms/step - loss: 1.3546 - accuracy: 0.5325 - val_loss: 1.7061 - val_accuracy: 0.4779
Epoch 38/100
19/19 [=====] - 3s 175ms/step - loss: 1.2781 - accuracy: 0.5634 - val_loss: 3.0292 - val_accuracy: 0.3382
Epoch 39/100
19/19 [=====] - 3s 182ms/step - loss: 1.2352 - accuracy: 0.5757 - val_loss: 1.4557 - val_accuracy: 0.5147
Epoch 40/100
19/19 [=====] - 4s 198ms/step - loss: 1.1643 - accuracy: 0.6096 - val_loss: 1.2538 - val_accuracy: 0.5588
Epoch 41/100
19/19 [=====] - 4s 185ms/step - loss: 1.2196 - accuracy: 0.5855 - val_loss: 1.5317 - val_accuracy: 0.5147
Epoch 42/100
```

```
19/19 [=====] - 3s 183ms/step - loss: 1.1232 - accuracy: 0.6201 - val_loss: 1.3263 - val_accuracy: 0.5147
Epoch 43/100
19/19 [=====] - 3s 182ms/step - loss: 1.1279 - accuracy: 0.6151 - val_loss: 1.6795 - val_accuracy: 0.4632
Epoch 44/100
19/19 [=====] - 3s 173ms/step - loss: 1.0944 - accuracy: 0.6054 - val_loss: 1.9967 - val_accuracy: 0.4485
Epoch 45/100
19/19 [=====] - 3s 184ms/step - loss: 1.2131 - accuracy: 0.5789 - val_loss: 3.8101 - val_accuracy: 0.2132
Epoch 46/100
19/19 [=====] - 3s 182ms/step - loss: 1.2451 - accuracy: 0.5773 - val_loss: 2.6766 - val_accuracy: 0.3897
Epoch 47/100
19/19 [=====] - 3s 174ms/step - loss: 1.2224 - accuracy: 0.5771 - val_loss: 1.6436 - val_accuracy: 0.4853
Epoch 48/100
19/19 [=====] - 3s 182ms/step - loss: 1.0597 - accuracy: 0.6266 - val_loss: 1.4739 - val_accuracy: 0.5147
Epoch 49/100
19/19 [=====] - 3s 167ms/step - loss: 1.1985 - accuracy: 0.5993 - val_loss: 3.4810 - val_accuracy: 0.3015
Epoch 50/100
19/19 [=====] - 3s 181ms/step - loss: 1.2855 - accuracy: 0.5724 - val_loss: 1.4647 - val_accuracy: 0.5515
Epoch 51/100
19/19 [=====] - 4s 196ms/step - loss: 1.1585 - accuracy: 0.6096 - val_loss: 1.6666 - val_accuracy: 0.5000
Epoch 52/100
19/19 [=====] - 3s 182ms/step - loss: 1.0927 - accuracy: 0.6151 - val_loss: 1.8896 - val_accuracy: 0.4559
Epoch 53/100
19/19 [=====] - 3s 175ms/step - loss: 1.1366 - accuracy: 0.6233 - val_loss: 2.2382 - val_accuracy: 0.3676
Epoch 54/100
19/19 [=====] - 3s 182ms/step - loss: 1.1186 - accuracy: 0.5921 - val_loss: 1.2781 - val_accuracy: 0.5956
Epoch 55/100
19/19 [=====] - 3s 168ms/step - loss: 1.3007 - accuracy: 0.5942 - val_loss: 2.1996 - val_accuracy: 0.4265
Epoch 56/100
19/19 [=====] - 3s 169ms/step - loss: 0.9481 - accuracy: 0.6610 - val_loss: 1.2351 - val_accuracy: 0.6029
Epoch 57/100
19/19 [=====] - 3s 174ms/step - loss: 0.9999 - accuracy: 0.6842 - val_loss: 1.5923 - val_accuracy: 0.4779
Epoch 58/100
19/19 [=====] - 3s 179ms/step - loss: 1.0877 - accuracy: 0.6053 - val_loss: 1.2925 - val_accuracy: 0.5662
Epoch 59/100
19/19 [=====] - 4s 202ms/step - loss: 1.0050 - accuracy: 0.6464 - val_loss: 1.3629 - val_accuracy: 0.5735
Epoch 60/100
19/19 [=====] - 3s 176ms/step - loss: 0.9799 - accuracy: 0.6644 - val_loss: 1.1559 - val_accuracy: 0.6176
Epoch 61/100
19/19 [=====] - 3s 169ms/step - loss: 0.9721 - accuracy: 0.6592 - val_loss: 2.2908 - val_accuracy: 0.3676
```

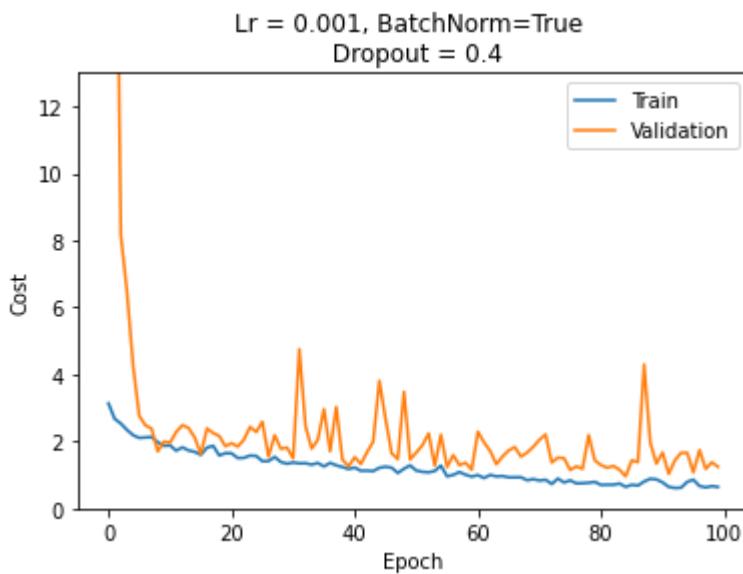
```
Epoch 62/100
19/19 [=====] - 3s 180ms/step - loss: 0.9081 - accuracy: 0.7023 - val_loss: 1.9711 - val_accuracy: 0.4485
Epoch 63/100
19/19 [=====] - 3s 168ms/step - loss: 1.0169 - accuracy: 0.6421 - val_loss: 1.7153 - val_accuracy: 0.5074
Epoch 64/100
19/19 [=====] - 3s 179ms/step - loss: 0.9569 - accuracy: 0.6760 - val_loss: 1.3210 - val_accuracy: 0.5882
Epoch 65/100
19/19 [=====] - 3s 184ms/step - loss: 0.9669 - accuracy: 0.6859 - val_loss: 1.5475 - val_accuracy: 0.4779
Epoch 66/100
19/19 [=====] - 3s 181ms/step - loss: 0.9368 - accuracy: 0.6849 - val_loss: 1.7305 - val_accuracy: 0.4559
Epoch 67/100
19/19 [=====] - 4s 185ms/step - loss: 0.9283 - accuracy: 0.6711 - val_loss: 1.8319 - val_accuracy: 0.4485
Epoch 68/100
19/19 [=====] - 3s 179ms/step - loss: 0.9256 - accuracy: 0.7038 - val_loss: 1.5465 - val_accuracy: 0.5000
Epoch 69/100
19/19 [=====] - 4s 192ms/step - loss: 0.8423 - accuracy: 0.7072 - val_loss: 1.6679 - val_accuracy: 0.5000
Epoch 70/100
19/19 [=====] - 3s 183ms/step - loss: 0.8748 - accuracy: 0.6990 - val_loss: 1.8416 - val_accuracy: 0.4779
Epoch 71/100
19/19 [=====] - 3s 172ms/step - loss: 0.8596 - accuracy: 0.7072 - val_loss: 2.0492 - val_accuracy: 0.5074
Epoch 72/100
19/19 [=====] - 3s 183ms/step - loss: 0.8557 - accuracy: 0.7039 - val_loss: 2.2084 - val_accuracy: 0.4118
Epoch 73/100
19/19 [=====] - 3s 168ms/step - loss: 0.7537 - accuracy: 0.7534 - val_loss: 1.3678 - val_accuracy: 0.5221
Epoch 74/100
19/19 [=====] - 3s 176ms/step - loss: 0.8908 - accuracy: 0.6941 - val_loss: 1.5174 - val_accuracy: 0.5735
Epoch 75/100
19/19 [=====] - 3s 179ms/step - loss: 0.7747 - accuracy: 0.7401 - val_loss: 1.5004 - val_accuracy: 0.5735
Epoch 76/100
19/19 [=====] - 3s 176ms/step - loss: 0.8416 - accuracy: 0.7106 - val_loss: 1.1474 - val_accuracy: 0.6397
Epoch 77/100
19/19 [=====] - 3s 173ms/step - loss: 0.7483 - accuracy: 0.7204 - val_loss: 1.2569 - val_accuracy: 0.6250
Epoch 78/100
19/19 [=====] - 3s 166ms/step - loss: 0.7368 - accuracy: 0.7380 - val_loss: 1.1814 - val_accuracy: 0.6324
Epoch 79/100
19/19 [=====] - 3s 177ms/step - loss: 0.8005 - accuracy: 0.7432 - val_loss: 2.1879 - val_accuracy: 0.4485
Epoch 80/100
19/19 [=====] - 3s 179ms/step - loss: 0.7923 - accuracy: 0.7303 - val_loss: 1.4326 - val_accuracy: 0.5809
Epoch 81/100
19/19 [=====] - 3s 179ms/step - loss: 0.6940 - accuracy: 0.7534 - val_loss: 1.5004 - val_accuracy: 0.5735
```

```
acy: 0.7620 - val_loss: 1.2920 - val_accuracy: 0.6029
Epoch 82/100
19/19 [=====] - 3s 182ms/step - loss: 0.7084 - accuracy: 0.7467 - val_loss: 1.2149 - val_accuracy: 0.6397
Epoch 83/100
19/19 [=====] - 3s 173ms/step - loss: 0.7390 - accuracy: 0.7723 - val_loss: 1.2683 - val_accuracy: 0.5882
Epoch 84/100
19/19 [=====] - 3s 172ms/step - loss: 0.7429 - accuracy: 0.7237 - val_loss: 1.1573 - val_accuracy: 0.6103
Epoch 85/100
19/19 [=====] - 3s 181ms/step - loss: 0.6451 - accuracy: 0.7829 - val_loss: 0.9571 - val_accuracy: 0.6618
Epoch 86/100
19/19 [=====] - 3s 173ms/step - loss: 0.6972 - accuracy: 0.7637 - val_loss: 1.4342 - val_accuracy: 0.5956
Epoch 87/100
19/19 [=====] - 4s 194ms/step - loss: 0.6841 - accuracy: 0.7829 - val_loss: 1.3808 - val_accuracy: 0.5368
Epoch 88/100
19/19 [=====] - 3s 183ms/step - loss: 0.8232 - accuracy: 0.7260 - val_loss: 4.2957 - val_accuracy: 0.2647
Epoch 89/100
19/19 [=====] - 3s 173ms/step - loss: 0.9306 - accuracy: 0.7158 - val_loss: 1.9639 - val_accuracy: 0.4191
Epoch 90/100
19/19 [=====] - 3s 175ms/step - loss: 0.8674 - accuracy: 0.7056 - val_loss: 1.3379 - val_accuracy: 0.6029
Epoch 91/100
19/19 [=====] - 3s 178ms/step - loss: 0.7736 - accuracy: 0.7312 - val_loss: 1.6693 - val_accuracy: 0.5147
Epoch 92/100
19/19 [=====] - 3s 183ms/step - loss: 0.6425 - accuracy: 0.7681 - val_loss: 1.0308 - val_accuracy: 0.6912
Epoch 93/100
19/19 [=====] - 3s 179ms/step - loss: 0.6113 - accuracy: 0.7961 - val_loss: 1.4377 - val_accuracy: 0.6029
Epoch 94/100
19/19 [=====] - 3s 175ms/step - loss: 0.6179 - accuracy: 0.7895 - val_loss: 1.6587 - val_accuracy: 0.5294
Epoch 95/100
19/19 [=====] - 4s 189ms/step - loss: 0.8173 - accuracy: 0.7106 - val_loss: 1.6605 - val_accuracy: 0.5662
Epoch 96/100
19/19 [=====] - 3s 179ms/step - loss: 0.8470 - accuracy: 0.7106 - val_loss: 1.0752 - val_accuracy: 0.6691
Epoch 97/100
19/19 [=====] - 4s 185ms/step - loss: 0.6691 - accuracy: 0.7878 - val_loss: 1.7435 - val_accuracy: 0.5588
Epoch 98/100
19/19 [=====] - 3s 177ms/step - loss: 0.6627 - accuracy: 0.7945 - val_loss: 1.1784 - val_accuracy: 0.6103
Epoch 99/100
19/19 [=====] - 4s 187ms/step - loss: 0.6669 - accuracy: 0.7730 - val_loss: 1.3755 - val_accuracy: 0.6103
Epoch 100/100
```

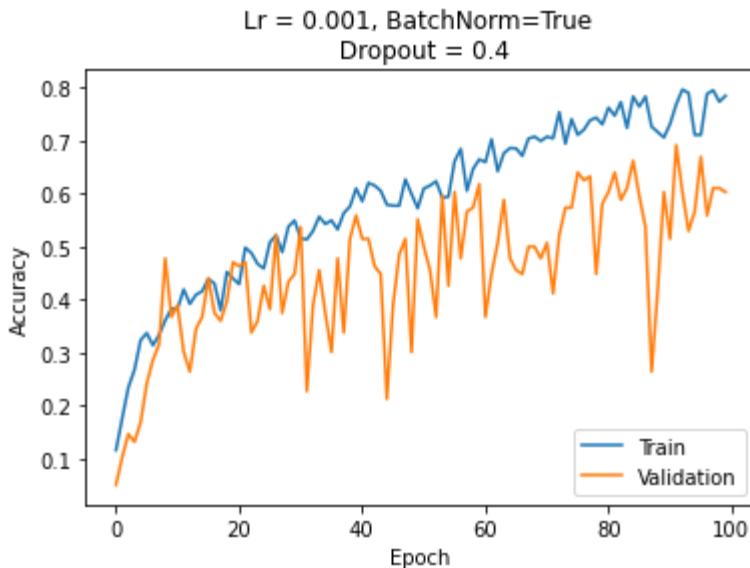
- Note: if you run `fit()` again, the model will continue training,

starting with the parameters it has already learnt, instead of reinitializing them.

```
In [45]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Lr = 0.001, BatchNorm=True \n Dropout = 0.4')
plt.ylabel('Cost')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper right')
plt.ylim(top=13)    # The instruction is used to limit the upper value of the
plt.ylim(bottom=0)  # The instruction is used to limit the lower value of the
plt.show()
```



```
In [46]: plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Lr = 0.001, BatchNorm=True \n Dropout = 0.4')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='lower right')
plt.show()
```



```
In [47]: # Predicting the image associated to the each sample in the test set (X_test)
predictions = oxford17_model.predict(test_x)
```

```
In [48]: print(type(predictions))
print(predictions.shape)
```

```
<class 'numpy.ndarray'>
(136, 17)
```

```
In [49]: # Predicting the image associated to the sample
# np.argmax returns the index of the maximum value
sample = 17
prediction = np.argmax(predictions[sample])
print("Prediction number=", prediction, ', it corresponds to a', dic[predicti
```

```
# Plotting the content of a sample
```

```
plt.imshow(train_x[sample]);
print('\ny =', np.squeeze(train_y[sample]))
```

```
for i in [i for i,x in enumerate(train_y[sample]) if x == 1]:
    print('')
```

```
print('y =', i, ';', 'the sample', sample, 'corresponds to a(an)', dic[i])
```

```
Prediction number= 3 , it corresponds to a ColtsFoot
```

```
y = [0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
```

```
y = 3 ; the sample 17 corresponds to a(an) ColtsFoot
```



In []: